



DARTMOUTH



Randomization & Approximation in Machine Learning and Neural Network Pruning

Math 76 Final Project

Warren Shepard Sri Korandla Tushar Aggarwal

August 2024

Dr. Alice Schwarze

Table of contents

1. The basics
2. k Approximate Nearest Neighbors
3. Graph-Based Semi-Supervised Learning
4. Randomized Pruning
5. Questions

Basic Randomization & Approximation in ML

Randomization & Approximation as we know it

The randomization techniques from class mostly involve **splitting** or **modifying data** randomly:

- splitting data into training and testing subsets

```
1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

- bagging/bootstrapping procedures
- random parameter/weight initialization in neural networks

Approximation is of course the goal of machine learning, but we often use *exact* algorithms in ML models instead of approximation algorithms

What else can Randomization & Approximation be used for

- Optimization:
 - trade "accuracy" for speed
 - oftentimes algorithms have $\geq 90\%$ accuracy
- In neural networks
 - "Prune" connections (somewhat) at random to get a smaller model with higher accuracy
- Sometimes, allowing some error is more "realistic" leading to a higher accuracy
 - higher accuracy's could also be because of a reduction in over fitting

Example: k Approximate Nearest Neighbors

Naive Solution


Problem: Given a set of vectors S in a data space and a similarity measure DIST (for example euclidean distance) find the k nearest ("most similar") neighbors to some input x .

Brute force solution: compare to every other point


```
1: procedure BRUTE-KNN( $x, S, \text{DIST}, k = 1$ )  
2:   Compute  $\text{DIST}(x, s) \forall s \in S$   
3:   return  $k$  points ( $s \in S$ ) with smallest  $\text{DIST}$  values
```

Issue: expensive and **redundant** (especially in higher dimensions and when there is lots of training data); $O(n)$ lookup


k Approximate Nearest Neighbors (kANN)



(A)



(B)



(C)

A) data points; B) & C) approximate nearest neighbor [3]

Locality Sensitive Hashing (LSH)


Hashing:

- hash function: $f: U \rightarrow [n]$ where $|U| \gg n$ and n is (usually) prime
- typically, two inputs have a very low probability of "colliding" (their output is the same)
 - inputs will have different "hashes" even if they are very similar

Locality Sensitive Hashing (LSH)


- two inputs which are similar have a high probability of colliding [3]
- can be used to solve kANN with high accuracy and in less time

LSH Preprocessing



Pre processing for LSH; completed in $O(n)$ time

LSH Query



Querying for LSH; completed in $O(1)$ time per query



Approximate-KNN Experiment

- using PyNNdescent package for kANN implementation [5]
 - uses modern kANN algorithm [1], written in python
 - accuracy converges to 90%
- Naive brute force implementation hand written in python
 - did not use an implementation from scikitlearn because compiled in C
 - best to have implementation we are comparing all in the same language

Results

```
# Generate a synthetic dataset with 10,000 samples and 1,000 features
```


```
X, y = make_classification(n_samples=10000, n_features=200, n_informative=50, n_classes=3,  
                           random_state=42)
```




kNN vs kANN **raw** accuracy and time comparison

Graph-Based Semi-Supervised Learning


**1) Start with data in feature space;
small percentage is labeled**




2) Construct k nearest neighbor Graph (here $k=2$)




3) "Propagate Labels"




3) "Propagate Labels"



3) "Propagate Labels"




Graph-Based Semi-Supervised Learning




Graph-Based Semi-Supervised Learning


1) Start with data in feature space;
small percentage is labeled




2) Construct k nearest neighbor
Graph (here $k=2$)




3) "Propagate Labels"




3) "Propagate Labels"



3) "Propagate Labels"



Complete!



Graph-Based Semi-Supervised Learning

- Constructing the k Nearest Neighbor Graph dominates runtime
- Most Label Propagation Algorithms (LPA) are **simple** and **near linear** in time

Parameters:

- $y :=$ initial labels; $y_i :=$ label of node i ; -1 if unlabeled
- $W :=$ adjacency matrix of nearest neighbors graph
- $\alpha \in [0, 1] :=$ "clamping factor" indicating influence of initial labels **per iteration**


Update Rules:

$$\bullet F_{ic}^{(0)} \leftarrow \begin{cases} 1 & \text{if } y_i = c \\ 0 & \text{otherwise} \end{cases} \quad F_{ic} \text{ is the } \mathbf{confidence} \text{ of node } i \text{ having label } c$$


$$\bullet F_{ic}^{(t+1)} \leftarrow \underbrace{\alpha W F^{(t)}}_{\text{propagate from neighbors}} + \underbrace{(1 - \alpha) F^{(0)}}_{\text{account for initial labels}}$$

$$\bullet y_i^* \leftarrow \arg \max_c F_{ic}^{(t_{\max})} \quad y_i^* \text{ is the vector of } \mathbf{final} \text{ labels}$$

Results I




Results I




Nearest Neighbor graph with $k = 5$

Results I




Final labels using $\alpha = .99$

Results II




- Time: 4.20 seconds
- Accuracy: 0.976





- Time: 2.89 seconds
- Accuracy: 0.977

MNIST Dataset Preprocessing



- Normalized signal to take value $\in [0, 1]$
- PCA to extract top 50 principle components
- masked 90% of labels to simulate semi-supervised learning; kept full label set for accuracy testing



MNIST Dataset Results


* Ran with A100 GPU runtime on Google Colab

* Parameters:

Using Brute Force kNN:

- Accuracy: 0.933
- Total Time: 18.81 minutes
- Time to construct kNN-graph: 17.14 minutes

Using k Approximate NN:

- Accuracy: 0.955
- Time: 1.79 minutes
- Time to construct kANN-graph: 

MNIST Dataset Results

- * Ran with A100 GPU runtime on Google Colab
- * Parameters:

Using Brute Force kNN:

- Accuracy: 0.933
- Total Time: 18.81 minutes
- Time to construct kNN-graph: 17.14 minutes

Using k Approximate NN:

- Accuracy: 0.955
- Time: 1.79 minutes
- Time to construct kANN-graph: < **6.2 seconds**

MNIST Dataset Results

* Ran with A100 GPU runtime on Google Colab

* Parameters:

Using Brute Force kNN:

- Accuracy: 0.933
- Total Time: 18.81 minutes
- Time to construct kNN-graph: 17.14 minutes

Using k Approximate NN:

- Accuracy: 0.955
- Time: 1.79 minutes
- Time to construct kANN-graph: < 6.2 seconds

Approximate NN-graph construction is over 165 times faster than the naive

The Problem With (our) Label Propagation (implementation)

- In theory, LPA can be parallelized
 - instead of executing **sequentially**, tasks are executed **simultaneously** on multiple units (i.e. CPU/GPU)
- No early stopping
 - the number of iterations is a parameter
 - stop early if there is no/little change between two consecutive iterations
- ~~Not leveraging sparse matrices~~
 - MNIST PCA is not sparse!

LPA Implementation Improvements & Results

Improvements:

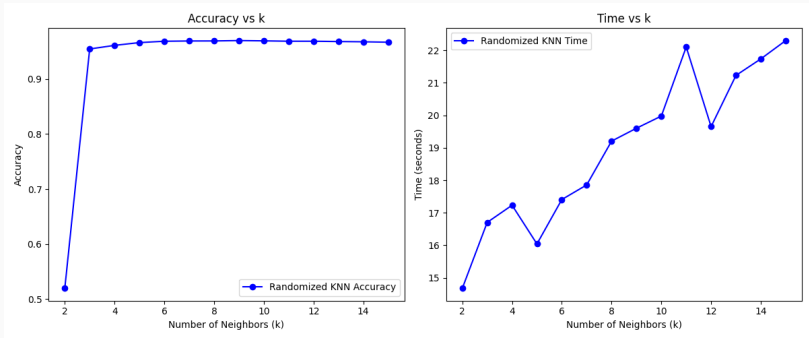
- Parallelization did not help
- Early stopping helped significantly and easy to implement

New results using k Approximate NN:

- Accuracy: 0.966
- Time: 16 seconds
- Time to construct kANN-graph: ≈ 6 seconds
- LPA time: **10 seconds; over 10 times faster**

Note: accuracy is 1% higher, which indicates initial model may have been slightly overfit

Hyper-Parameter (k) tuning




Accuracy (left) and time (right) comparison for $k \in [2 : 15]$

Note: Increase in time is a result of NN-graph construction, not LPA

The Inefficiency with LPA (in this context)

- we do not consider distances to nearest neighbors
- however, we already calculated these distances in kNN
- in kANN we did not, but it would take negligible time to compute distances
- Idea: if we consider distances to nearest neighbors in LPA, it could potentially be more accurate




Randomized Pruning in Neural Networks

Why Prune?

- "It is widely acknowledged that large and sparse models have higher accuracy than small and dense models" [4]
- pruning a model prevents over-fitting
- pruned models also require a fraction of the space a unpruned model would need in memory

Neurology Inspiration



Synaptic pruning in the brain [2]

Dataset and Neural Network Architecture

- We use a simple architecture recommended by tensorflow
- We prune connections on one layer (more detail on next slide)
- pruning is done by setting pruned weights (outgoing to next layer) to zero
- Inspiration for randomized pruning: "Breaking through Deterministic Barriers: Randomized Pruning Mask Generation and Selection" [4]



MNIST Fashion Dataset

RESULTS

Network	Pruning Strategy	Accuracy
No Pruning	No pruning is applied. All connections are retained.	87.68%
Threshold Fraction Pruning	Prune a specified fraction of connections with the smallest weights.	86.63%
Random Pruning	Prune a random subset of connections, regardless of their weights or importance.	64.56%
Importance-Weighted Random Pruning	Prune connections probabilistically based on their importance, where more important connections are less likely to be pruned.	78.06%
Fractional Random Threshold Pruning	First, select a fraction of connections with the smallest weights. Then, randomly prune a smaller fraction of these selected connections.	87.47%

Summary of Pruning Strategies for Five Neural Networks

Questions???



W. Dong, M. Charikar, and K. Li.
Efficient k-nearest neighbor graph construction for generic similarity measures.
[In *Proceedings of the 20th international conference on World wide web*, pages 577–586. ACM, 2011.](#)



J. Embrace.
Synaptic growth, synesthesia, and savant abilities.
[Embrace Autism, 2024.](#)
Accessed: 2024-08-12.



S. Har-Peled, P. Indyk, and R. Motwani.
Approximate nearest neighbor: Towards removing the curse of dimensionality.
[Theory of Computing, 8:321–350, 2012.](#)
Special issue in honor of Rajeev Motwani.



J. Li, W. Gao, Q. Lei, and D. Xu.

Breaking through deterministic barriers: Randomized pruning mask generation and selection.

[arXiv preprint arXiv:2310.13183](#), 2023.

Accessed: 2024-08-14.



PyNNDescent Developers.

[PyNNDescent Documentation](#), 2024.

Accessed: 2024-08-14.